DRILLLOOP / PUBLIC-SAFETY SECTOR

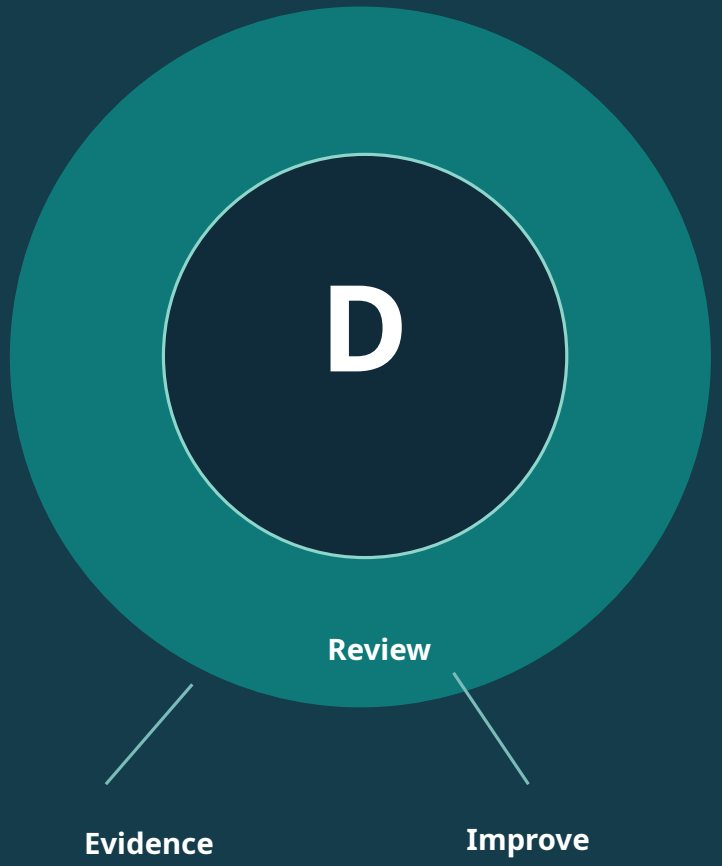
DrillLoop

Evidence-Linked Exercise Review and Action Closure

For hospital preparedness evaluators · synthetic demo

WORKING MVP

Fictional hospital prolonged power outage · 27 October 2026



The claim stays connected to its source and its reviewer decision.

THE PROBLEM

The exercise ends. The evidence work begins.

Evaluation is a reconciliation job across files, clocks, objectives, and people.

01

Evidence is scattered

Evaluator notes, control-room logs, transcripts, and reports use different formats and detail.

02

Timelines disagree

Missing times, conflicting statements, and recognition errors can make event order look more certain than it is.

03

Actions lose momentum

Findings become difficult to assign, review, and verify after the report is written.

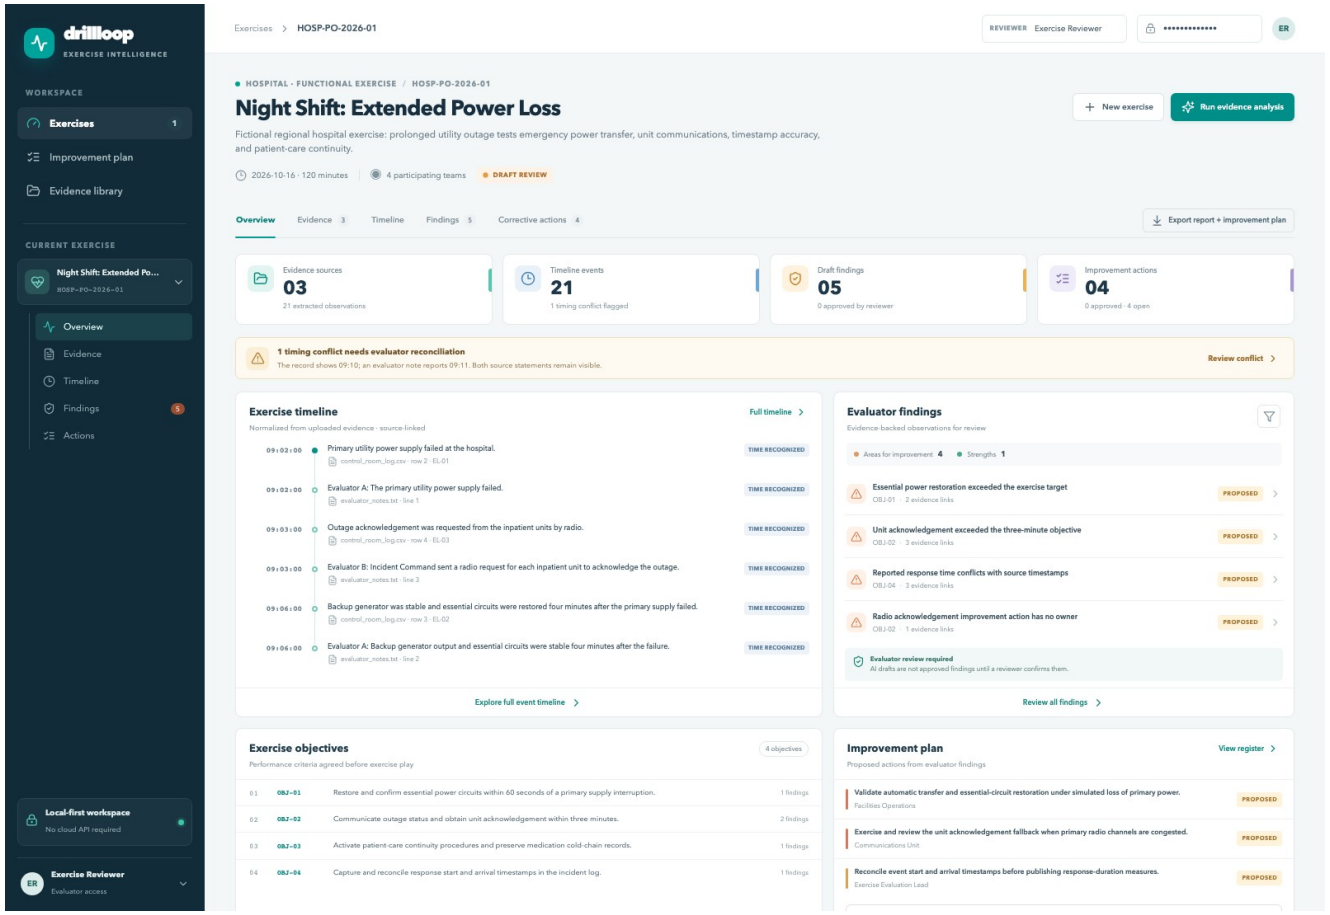
What teams need

A review trail that shows what the evidence says, what the evaluator concludes, and who verifies the improvement.

The amount of time saved and the size of the gap vary by organization; neither has been measured yet.

One review workspace for the whole exercise

Follow evidence through to an after-action report and improvement plan (AAR/IP).



Screenshot captured from the running app. All exercise content is fictional.

SYNTHETIC HOSPITAL EXERCISE

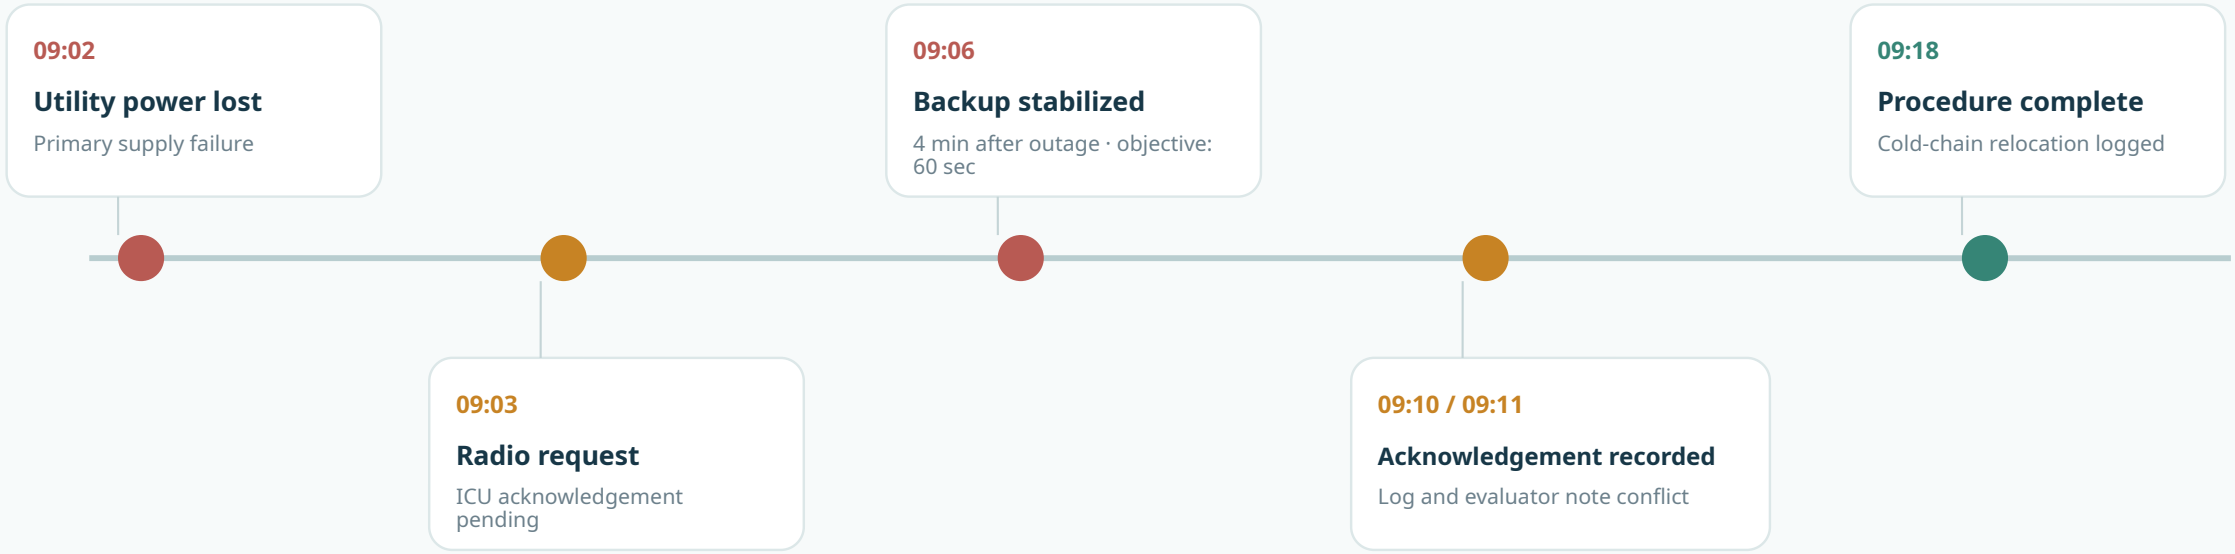
The screen keeps together:

- **Exercise plan**
  Objectives and milestones
- **Evidence**
  File and source locator
- **Timeline**
  Known, missing, and conflicting times
- **Review**
  Proposed vs approved status

DEMO SCENARIO

A power outage with evidence that does not all agree

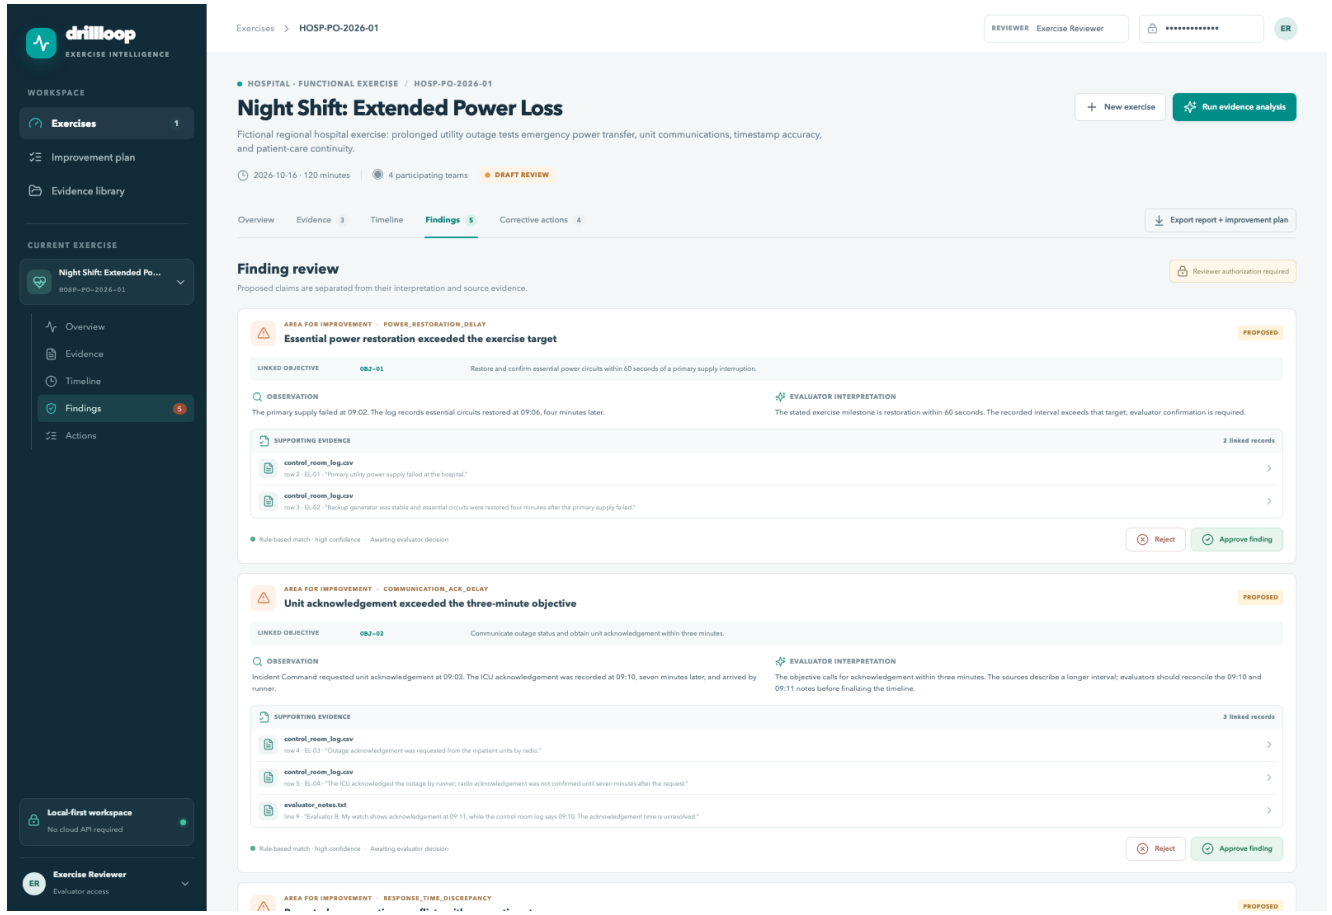
Hospital HOSP-PO-2026-01 · Night Shift: Extended Power Loss · fictional exercise.



Uncertainty stays visible 09:10 vs 09:11 is unresolved. “Generator lagged a bit” has no time or measure, so it is not ordered as a confirmed event.

Open the source behind every draft

See what the record says, what it may mean, and where a reviewer can check it.



No finding is approved by the model. Reviewer decision remains explicit.

DRAFT FINDING

Delayed ICU acknowledgement

WHAT IT MAY MEAN

About 7 minutes observed; exercise target is 3 minutes.

SUPPORTING RECORDS

CSV rows 4–5

Note line 9

DRAFT / awaiting evaluator

The response-time discrepancy cites the independent event log; the low-confidence “too minutes” transcript phrase remains uncertain.

AI WORKFLOW

A fixed pipeline keeps the evidence path inspectable

The MVP uses deterministic extraction for the demo fixture. Optional local models stay behind defined boundaries.

01

Collect files

Notes · spreadsheets
transcripts · readable PDFs

02

Record details

Keep exact file location
and known / unknown time

03

Compare

Compare recorded events
with exercise targets

04

Link evidence

Show the exact source
behind every proposal

05

Human review

A person decides
approve / reject

OPTIONAL LOCAL SPEECH-TO-TEXT

faster-whisper can turn audio into draft text when installed. Demo transcript is pre-written; speakers are not identified.

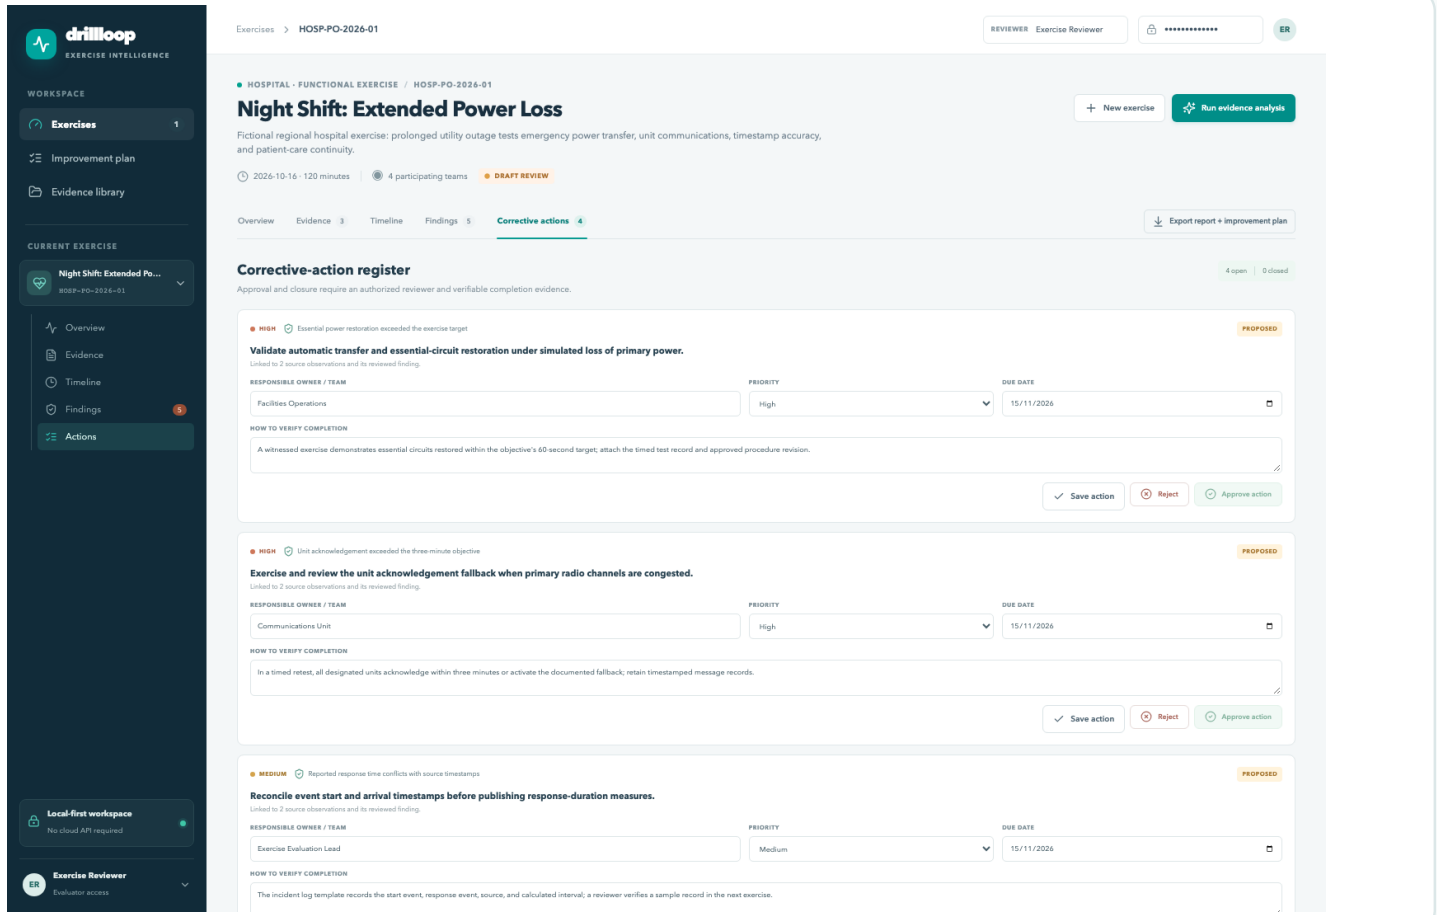
OPTIONAL LOCAL WRITING MODEL

Ollama may reword an action draft only. It cannot create a finding, change a source link, or make a reviewer decision.

Same evidence + same rules produce a repeatable draft. The evaluator still decides.

A finding stays open until its improvement is checked

Every action needs an owner, a due date, a clear test for completion, and reviewer confirmation.



Proposed
Finding awaiting review

Assigned
Owner + due date

In progress
Acceptance criteria

Closure review
Attach completion proof

Closed
Reviewer confirms

Missing owner stays unassigned.

MARKET CONTEXT

The category already has AAR and action-tracking tools

DrillLoop should validate a focused workflow advantage instead of claiming report generation is unique.

OPTION / CONTEXT	PUBLICLY VISIBLE SCOPE	IMPLICATION TO TEST
FEMA exercise guidance	Exercise planning and evaluation steps	Use familiar structure; no claim of certification.
Veoci	Publicly describes AAR reports, timestamps, transcripts, and corrective actions	Direct competitor; compare source-level review and closure in customer demos.
D4H	Emergency-team readiness, incident, personnel, and equipment workflows	Adjacent incumbent; confirm evaluation evidence workflow depth.
DrillLoop hypothesis	Exact source observation → human decision → closure evidence	Potential niche; uniqueness and willingness to pay remain unvalidated.

Sources: FEMA exercise guidance · Veoci after-action reporting · D4H team readiness. Public pages, not a scored teardown.

Start with a hospital exercise review

This is a proposed first use case; no pilot partner or buyer has been confirmed.

FIRST PILOT HYPOTHESIS

Hospital tabletop review

Initial segment hypothesis. The included outage exercise is synthetic; no hospital partner is confirmed.

Report preparer

Exercise evaluator / preparedness staff (candidate user).

Report approver

Preparedness lead or exercise sponsor (candidate).

Budget owner

Not identified; role to validate in discovery.

Roles to validate

1. Who prepares the AAR/IP?
2. Who approves findings?
3. Who controls the budget?

No interviews or purchasing-path validation are recorded yet.

Start with a per-exercise pilot fee

This is the first model to validate; no price, buyer demand, or revenue has been established.

FIRST MODEL TO TEST

Per-exercise pilot fee

Price to be set with the prospective buyer after delivery effort and procurement path are understood.

BUYER — **Who owns the budget and approves an exercise-level purchase?**

VALUE — **Does evidence lookup, review, or action follow-up improve on the current process?**

DELIVERY — **What setup, security, support, and evidence-handling work does each exercise require?**

Annual plans and private deployment remain later hypotheses.

TECHNICAL FOUNDATION

Simple local app, with a person at the review gate

The demo can run without a cloud AI service; optional models stay on the computer.

Website screen

Exercise setup
Evidence review
Action tracking
React + TypeScript

Application service

Checks uploaded files
Applies review rules
FastAPI + Pydantic

Saved records

SQLite for the demo
PostgreSQL for deployment
SQLAlchemy database tools

Reports + optional local AI

PDF report: ReportLab
Speech-to-text: faster-whisper
Action wording: Ollama

TRACEABLE RECORD

Original file + file-change checksum + extracted quote + location + reviewer decision + completion proof

DEMO SECURITY BOUNDARY

Shared demo key; no named accounts, tenant separation, retention controls, or security assessment.

Exact library versions and licenses are listed in the Technical Design.

TECHNICAL EVALUATION

The sample passed its checks; new exercises still need testing

Scores use one fictional exercise built for this demo. They do not measure evaluator agreement or general accuracy.

EVENT RECORDS

8 / 8

All labeled rows found

FINDING LABELS

5 / 5

All expected labels found

SOURCE LINKS

9 / 9

Expected record IDs linked

RECORDED TIMES

8 / 8

One disagreement shown

Measured

Rule request latency p50 4.61 ms · p95 9.13 ms
Peak process RSS 78.2 MB · no model loaded

Not measured

Unseen-exercise errors · reviewer edits · semantic citation quality
ASR WER/latency · local LLM quality and latency

Precision and recall were 100% because rules and labels used the same small scenario; they do not establish performance on an unfamiliar exercise.

31 / 32 action fields were filled; one owner is intentionally missing. No evaluator study was run.

12 local test calls · no AI model loaded · results apply to this synthetic fixture only.

THE NEXT PROOF POINT

Pilot the review trail.

Compare the current process and DrillLoop on the same permitted hospital exercise records, then test an unfamiliar exercise.

2–3 users

Prospective hospital-preparedness interviews (proposed)

Paired review

Same permitted records: existing process and DrillLoop

Measure

Time · misses · corrections · citations · action fields

Partner-defined success thresholds.

No pilot partner, interviews, baseline, or performance targets are confirmed.